

The Anchor Does Most of the Work

◆ The verdict

Negotiation is a real, trainable skill — but the peer-reviewed evidence supports far fewer moves than the bookshelf implies, and the strongest one is almost boring: **whoever puts a credible number on the table first usually sets the price.** The first offer explains roughly **50–85% of the variance** in the final outcome (~90% confidence). The protocol isn't universal — it *inverts* between one-shot distributive deals (anchor hard, precise, aggressive) and repeated relational ones (warmth, perspective-taking, never anger) — but **asymmetrically**: aggression and anger flip across contexts, while warmth and cognitive perspective-taking help in *both* (~80% confidence). The decision that dominates every tactic is informational, not psychological: if you know their reservation value, anchor first; if you don't, let them open or ask diagnostic questions. The famous books are directionally right, memorable, and oversold on scope — **Chris Voss's exact scripts have never been tested in a controlled negotiation**, and Fisher–Ury's win-win universalism breaks in pure-distributive and cross-cultural fights (~85% confidence). The popular *summaries* are worse: the viral "**38% more**" salary statistic traces to nothing, "always use precise numbers" reverses with experts, and "women don't ask" is now contested by the field that launched it. Where I'd put my chips: learn the six lab-validated levers, diagnose the deal, and treat the books as a phrasing layer — not as evidence.

◆ The one number that runs the field

Everything downstream rests on one finding. The keystone is **Petrowsky, Schweinsberg et al. (2025), "The power and peril of first offers,"** in *Organizational Behavior and Human Decision Processes* — a preregistered synthesis of **374 effect sizes across 90 samples and 16,334 participants**, plus replications with another **2,121**. It quantifies both sides of the anchor at once:

Effect	Size (Hedges' g)	Effect sizes (m)	What it means
First offer → agreement value	r = 0.62 / g = 1.56	53	The opening number all but is the final number
First-mover advantage	g = 0.42	80	Going first beats going second, on average
Ambitious vs. moderate open	g = 1.14	187	Aiming high pays — when a deal closes
More impasses (the cost)	g = -0.42	13	Aggression breaks more deals
Worse counterpart satisfaction	g = -0.40	41	And sours the ones that survive


The first offer correlates with the settlement at $r = 0.62$. In distributive terms, the person who speaks the first credible number is doing most of the negotiating — the rest is adjustment around their anchor.

The mechanism split is the whole game in one sentence: **selective accessibility drives the counteroffer** (their mind fills with anchor-consistent facts and adjusts too little), while **anger drives the impasses and the resentment**. That is *why* the protocol inverts. The anchor's cognitive benefit is nearly universal; its emotional cost is what wrecks any deal you have to live inside afterward.

Two honest caveats the pop summaries skip. First, the "peril" effects rest on far thinner evidence — **13 and 41 effect sizes** versus 80–187 for the "power" side — so the downside is real but more loosely estimated than the upside. Second, complexity attenuates everything: *"as the number and type of issues increased, the effects of first offers became smaller, and the mechanisms changed."* The anchor is a scalpel for the single-issue haggle and a blunter instrument for the ten-variable partnership.

◆ The decision that comes before every tactic

Before "what do I say" comes a prior question that decides more than any script: **do you know their reservation value or the market range?** It's the most consequential move in the discipline, and it has a clean operating rule from Galinsky's "Should You Make the First Offer?" (HBS/PON, 2004) and the 2025 meta-analysis.



When you know the range (anchor first). Make an aggressive-but-justifiable opening offer. "Justifiable" is the constraint doing the work: Malhotra and Bazerman's rule is the most extreme number you can defend with a standard *without embarrassment*. In salary, where compensation bands are now Googleable on levels.fyi and Glassdoor, you almost always know the range — which is exactly why the folk rule "never make the first offer" is worst precisely where most people apply it.

When you don't know the range (let them open, or ask first). This is the real content behind Voss's "let the counterpart anchor" default — you can't anchor from ignorance without either leaving money on the table or triggering impasse, and their offer is itself free information. Loschelder, Trötschel, Swaab & Galinsky (2014, *Psychological Science*) named the **First-Mover Disadvantage**: when your offer leaks compatible preferences, you claimed only ~44% of available profit. Under information asymmetry — you the layperson versus a market expert — the first-mover advantage can *reverse* (Loschelder et al. 2016), with informed recipients capturing ~71% vs. ~52% of profit. If you're the amateur at the pawnshop, do not open.

The bridge between the branches is the diagnostic question — the move that converts "unknown" into "known" before anyone anchors. A field experiment in *Negotiation Journal* (2023) found that the low-commitment opener **"What is your best price?"** extracts a real concession and surfaces the seller's floor *without you anchoring yourself*. And whoever opens, the same debiasing protects you from *their* anchor: Galinsky & Mussweiler (2001) showed that shifting attention onto your own target and their likely walk-away — rather than ruminating on their number — blunts the anchor. Counter-anchoring *attenuates* the first-mover advantage; it does not *eliminate* it in the pooled data, where $g = 0.42$ survives.

◆ The six levers, weighed

Strip away the branding and the field supports a short list. Here is each lever with its effect size, its home turf, and the condition under which it turns on you.

Lever	Evidence	Works best when	Backfires / inverts when	Anchor source
First offer / anchor	$r = 0.62$; ~50–85% of variance	You know the range; single-issue	You're the uninformed party; you leak compatible info	Galinsky & Mussweiler 2001; Petrowsky 2025
Precise number	Smaller counter-adjustments; field-confirmed	Justified, realistic precision	Fake precision with experts (inverted-U)	Mason et al. 2013; Loschelder 2016
Bolstering range	Beats point offers, ~2× fewer impasses than a lone high number	You want to claim <i>and</i> stay likable	You use a "bracketing/backdown" range instead	Ames & Mason 2015 (5 studies)
Avoid anger	Anger raises concessions but lowers trust/value (64-study meta)	Never a first resort	Repeated games, strong-BATNA targets, faked anger	Van Kleef 2004; Sharma et al. 2020
Positive affect	Gains concessions in distributive <i>and</i> integrative	Always — it's near-free	(No documented inversion)	Kopelman et al. 2006; Carnevale & Isen 1986
Perspective-taking	Creates <i>and</i> claims value	Always — cognitive, not emotional	You substitute <i>emotional empathy</i> → over-concede	Galinsky, Maddux et al. 2008

The deepest insight in the whole literature — and the correction to the "distributive vs. integrative" binary everyone teaches — is that the inversion is **asymmetric**. Two levers genuinely flip. Two do not.

	Distributive one-shot	Integrative repeated
▲ Anchor aggressively	DO	AVOID
▲ Show anger	RISKY EDGE	AVOID
● Warmth / positive affect	DO	DO
● Perspective-taking	DO	DO
▲ inverts ● holds in both		

The practical translation is liberating: **you can anchor hard and be warm**. The pop-advice error is bundling "assertive" with "cold" — putting aggression in the *manner* when it belongs only in the *number*. Kopelman, Rosette & Thompson's "Three Faces of Eve" (2006) found positive emotional display won *more* concessions even in a distributive setting; Galinsky et al. (2008) found cognitive perspective-taking (modeling what they want and why) creates and claims value, while *emotional empathy* leads to unnecessary self-sacrifice. So the sharpest reframe of Voss's celebrated "tactical empathy" is that it works exactly to the extent it is cognitive perspective-taking — and the word "empathy" points the wrong way.

◆ The protocols, with exact phrasings

This is what you actually need at the table. Confidence tags: **[V] = lab-validated construct**, **[M] = mechanism-plausible but not outcome-tested** (i.e., the Voss scripts).

Salary — the canonical hybrid (distributive number, relational aftermath)

Both mechanisms fire here, which is why naïve "anchor as high as you can" advice is unsafe. Do this:

- **Convert unknown → known.** Pull the band from levels.fyi, Glassdoor, Pew, and internal comps. Now you're in the "anchor first" branch.
- **Anchor precise, justified, at the top of the credible range — not the sky.** Say: "*Based on my research, comparable roles at [peer firms] are landing between \$150K and \$165K. Given [specific value you bring], I'm targeting \$158,000.*" **[V]** Precise + justified + upper-reasonable respects Maaravi et al.'s (2022) finding that the anchoring-response function *flattens* at extreme demands and high asks depress likability and hiring odds.
- **Or open with a bolstering range:** "*I'm looking for \$155K–\$170K*" — your real target (\$155K) at the floor. **[V]** Ames & Mason (2015) found this claims more *and* reads as more cooperative; a lone high number roughly **doubles the impasse rate**.
- **Never anchor on your current salary or a weak alternative.** Hold your *target*. (See the weak-BATNA paradox below.)

- **If you hit backlash**, buffer with a relational account: *"I'm genuinely excited about the team — I want us to land somewhere that reflects the market for this role."* [V, partial] Bowles' work suggests signaling communal intent softens the penalty, though this is the identity-non-neutral part of the game.

Purchase — the near-pure distributive case (car, house, marketplace)

- **Don't know their floor?** Let them anchor, or ask *"What's your best price?"* [V] — extract the floor without anchoring yourself.
- **Know the market?** Anchor first, precise and justified against comparables.
- **The Ackerman gambit** (Voss, [M]): set your ceiling, open at **65%** of it, then climb **85% → 95% → 100%** in *shrinking* increments (the decreasing steps signal you're near your limit); make the final number oddly precise (**\$37,893**, not \$38,000); close with a small non-monetary throw-in.
- **Resist their sticker:** *"How did you arrive at that number?"* [M/V] buys time and forces justification while you re-anchor on your target and walk-away, not their figure.

Conflict — the relational case (coworker, partner, ongoing dispute)

Here the anchor is a trap and anger is a slow-acting poison. Flip the whole stance:

- **Do not open aggressive. Do not deploy anger** — the retaliation and trust-loss spill into every future interaction.
- **Perspective-take out loud:** *"What would need to be true for this to work on your end?"* [V] surfaces interests and hidden trades (logrolling).
- **Label the emotion to cool it:** *"It seems like this deadline is putting you in a real bind."* [M] Affect-labeling has neuroscience grounding (Lieberman et al. 2007 — naming a feeling dampens amygdala reactivity), but was never tested as a negotiation-outcome RCT.
- **Anchor on a shared standard, not a position:** *"What's a fair way for both of us to decide this?"* [V] — Fisher–Ury's "objective criteria," vindicated by the fixed-pie-bias research (De Dreu, Koole & Steinel 2000) showing negotiators chronically miss integrative trades.

The phrasebook

Situation	Say this	Why it works	Grade
Open, informed	"Comps land at \$X–\$Y; I'm targeting \$158,000 ."	Precise + justified anchor	[V]
Open, claim + stay liked	"I'm looking for \$155K–\$170K ." (target at floor)	Bolstering range	[V]
Open, uninformed	"What's your best price?"	Extracts floor, no self-anchor	[V]
Resist their anchor	"How did you arrive at that number?"	Forces justification; refocuses off anchor	[M/V]
Decline without conflict	"How am I supposed to do that?"	Calibrated "no"; hands them your problem	[M]
Cool a hot moment	"It seems like this is really stressful for you."	Affect labeling	[M]
Lock in agreement	Get them to say "That's right."	Felt understanding	[M]
Find trades	"What would need to be true for this to work?"	Surfaces interests → logrolling	[V]

◆ What the famous books get wrong

The books are not *wrong* about whether their levers work — they're wrong about **scope**, which is precisely where anecdote is weakest and meta-analysis strongest. When a book says "always" or "never" and the lab shows a moderator, the moderator wins, because someone measured it.

Popular claim	Verdict	What the research says
"Never make the first offer"	Conditional, not a law	Right only under information disadvantage; forfeits the 50–85% anchoring lever when you're informed
"Split the difference is fair" (Voss's title)	Right in distributive haggling	Reflexive compromise leaks value the anchor already captured; weak point in integrative logrolling
"Women don't ask — that's the pay gap" (<i>Women Don't Ask</i>)	Now contested	Artz/Goodall/Oswald (2016): women ask as often as men; cause shifts to backlash + structure
"Interests, not positions, solves everything" (<i>Getting to Yes</i>)	Overreach	Many deals are irreducibly distributive; the first <i>position</i> still drives the number ($r = 0.62$)
"Tactical empathy" wins	Right for the wrong reason	It works as <i>cognitive</i> perspective-taking; <i>emotional</i> empathy makes you over-concede
Voss's scripts work everywhere	Genre mismatch	Derived from hostage crises: extreme power asymmetry, non-monetary, one-shot, no BATNA for the taker

Two book-specific reckonings:

Never Split the Difference (Voss, 2016) is a set of memorable phrasings with plausible mechanisms, oversold as a validated system. I searched for peer-reviewed outcome tests of labeling, mirroring, "That's right," accusation audits, and calibrated questions — there are none. The *adjacent* constructs have support, but even that is a loose match: Maddux, Mullen & Galinsky (2008) found **67% of dyads reached agreement when one side subtly mimicked the other, versus 12.5% without** — but that's *postural/mannerism* mimicry, not Voss's "repeat their last three words." This channel has met the trap before — *How Old Is the World's Oldest Fire? Nobody Has Actually Dated It*: a claim repeated for years, taught in seminars, quoted as settled, and never once *measured*. Voss's scripts are the world's-oldest-fire of negotiation advice — probably useful, genuinely famous, undated.

Not one of Voss's branded techniques — labeling, mirroring, calibrated questions, "That's right" — has a controlled negotiation-outcome study behind it. The constructs beneath them do; the scripts do not.

Getting to Yes (Fisher & Ury, 1981) earned a takedown from *inside* the academy: James J. White's 1984 review in the *Journal of Legal Education* charged it with relying on anecdote and denying the reality of hard distributive conflict, and McKeown's later critique flags its cultural boundedness (bottom-line-first cultures break the model). Its lasting contribution — BATNA — is real and empirically robust, but the book tells you to *have* one, not how to *improve* one.

◆ Five things the summaries get wrong — including one outright error

Even careful readers absorb corrupted versions of the research. Five fixes:

1. The "38% more" salary stat is a ghost. Trace it and there is no study showing negotiators capture 38% over the initial offer. The nearest real number is **Pew (2023)**: of workers who asked for more, **~28% got exactly what they requested, ~38% got some increase, and ~35% got nothing.** The 38% is a *share of people in a middle bucket*, not a *magnitude of raise* — a statistic that shed its meaning in transit. The companion "**~44% negotiate**" is directionally fine; Pew and CareerBuilder both put non-askers near 55–60%.

2. "Always use precise numbers" has a documented reversal. Precise anchors do work — Mason et al. (2013) — but **Loschelder et al. (2016), "The Too-Much-Precision Effect"** (5 experiments, 1,320 people) found an *inverted-U for experts*: excessive precision reads as *incompetence* and backfires unless a rationale restores perceived competence. The field resolves the tension in precision's favor when precision is *realistic*: Hukkanen & Keloharju (2019) found that in real M&A deals, round bids cost buyers money.

In real mergers, a median round-numbered offer ran roughly **\$4–5 million higher** than a precise one — and closed less often. Fake precision insults experts; *realistic* precision signals you did the valuation. (Observational data — careful bidders may both price precisely and negotiate better.)

3. Gender backlash is real but small and heavily moderated. Bowles, Babcock & Lai (2007) is genuine and replicated. But **Mazei et al. (2015)** — 123 effect sizes, ~10,900 people — found men's economic edge is *small* and *vanishes* when negotiators have experience, know the bargaining range, or negotiate on someone's behalf. Backlash is a real tax, not the whole pay gap, and it is largely dissolvable with information.

4. The weak-BATNA paradox breaks "always improve your alternative." Schaerer, Swaab & Galinsky (2015) found negotiators with *no* alternative made *higher* first offers and got *better* outcomes than those with a *weak* one — because a weak alternative self-anchors you downward, while having none frees your anchor. The fix isn't a better BATNA; it's focusing on your **target**, not your fallback.

5. Anger is a high-variance distributive tool, not a strategy. Van Kleef (2004) is the result pop advice loves — angry negotiators extract concessions. But **Sharma et al.'s (2020) 64-study meta-analysis** confirms every boundary condition: anger works only on low-power targets who can't retaliate and can't detect that it's faked. Against a strong BATNA it does nothing; in a repeated relationship it detonates later.

◆ How the thinking evolved

The field's history is a sixty-year migration from prescriptive frameworks toward measured effect sizes — and the newest data is finally closing the lab-versus-field gap that a skeptic would exploit.

1965 Dist. vs integrative

1974 Anchoring

1981 Getting to Yes

2001 First offers = anchors

2007 Gender backlash

2015 Bolstering range

2016 Voss · precision cap

2025 Meta: 374 effects

Walton & McKersie (1965) named the distributive/integrative split still taught today. Tversky & Kahneman (1974) supplied the cognitive engine — anchoring-and-adjustment — that would later explain first offers. Fisher & Ury (1981) gave us BATNA and "win-win," and Galinsky & Mussweiler (2001) turned anchoring into a negotiation prescription with real numbers. The 2010s were the moderator decade: backlash mapped and then *bounded* (Mazei 2015), the counterintuitive bolstering range (Ames & Mason 2015), the weak-BATNA paradox (Schaerer 2015). Then 2016 delivered the field's split personality — Voss's blockbuster of untested scripts and Loschelder's austere "precision has a ceiling." The 2025 Petrowsky meta-analysis, alongside field datasets of **26 million eBay negotiations** and real M&A bids, is the reconciliation: the lab effects largely survive contact with real markets.

◆ The strongest case against everything above

Let me steelman the opposite of my verdict — that **the books are right and the lab research is the thing that's oversold** — because it has real force (I hold my position at ~78%, and here is the 22%).

- **External validity is genuinely shaky.** The headline effects come disproportionately from student subjects, small stakes, one-shot artificial tasks. A $g = 0.42$ in a classroom simulation could wash out to near-zero in a real deal dominated by reputation, relationships, and repeated play.

"Replicated" often means *replicated within the same narrow paradigm* — replicated narrowness, not validated breadth.

- **The books were pressure-tested where it counts.** Voss's methods came from thousands of live crises where being wrong meant someone died. That's survivorship-biased, yes — but it's also selection for *what executes under stress*, which no lab captures.
- **Nuance is unusable at the table.** "It depends on the distributive/integrative ratio and your information about their reservation value, and remember the inversion is asymmetric" is unrunnable in a live, emotional negotiation. A crisp heuristic that's 70% right and *executable* beats a calibrated framework you freeze trying to apply. The books win on *usability*, which is the only metric that matters in the room.
- **The research's own caution is thin.** The meta-analysis's "peril" side rests on **13 effect sizes**. If the field's warning against aggression is itself loosely estimated, maybe the books' "aim high, go first" instinct is the better default and the caveats are noise.

The honest synthesis this forces: the books are the *executable layer* and the research is the *edge-case manual*. My rebuttal — and why I still land at ~78% for the research-first view — is that the lab-field gap is closing fast: real-estate agents were provably anchored by list price (Northcraft & Neale) and *denied it*; M&A precision is field data; the eBay corpora are 26 million real deals. The direction of the last decade's evidence is *toward* the lab, not away. But if you handed me a live, high-stakes, relationship-laden negotiation tomorrow with no prep time, I'd reach for the executable script first and the effect sizes second — and that concession is the 22%.

◆ What would change this answer

Concrete evidence or events that would move the verdict:

- **A field RCT of Voss's specific scripts** — labeling, mirroring, calibrated questions — showing a measurable outcome effect. This would upgrade the books from "plausible UI" to "validated," and I'd raise their status sharply. Its continued absence is why they stay graded **[M]**.
 - **The Petrowsky peril effects failing to replicate** at larger m . If impasse ($g = -0.42$, $m = 13$) and dissatisfaction ($g = -0.40$, $m = 41$) shrink toward zero with more studies, the "anger wrecks relational deals" pillar weakens and aggressive anchoring gets safer across the board.
 - **Field evidence that anchoring washes out** in genuinely high-stakes, repeated, professional negotiations. This is the steelman's strongest bet; the eBay/M&A data currently cuts the other way, but a well-powered null in elite dealmaking would move me toward the books' relationship-first view.
 - **Larger, less-moderated gender backlash in real hiring.** If field studies show backlash that *doesn't* dissolve with information and experience, the "small and moderated" correction reverses and the identity-neutrality of the playbook collapses.
 - **Cross-cultural data breaking the Western model.** Strong evidence that anchor-first, interests-based negotiation systematically fails in bottom-line-first or high-context cultures would demote the whole framework from "how negotiation works" to "how *Western* negotiation works."
-

◆ Sources, annotated

- [Petrowsky et al. 2025, "The power and peril of first offers" \(full PDF\)](#) — The keystone. 374 effect sizes, preregistered; every headline number verified against primary text. Trust: very high.
- [Galinsky & Mussweiler 2001, "First Offers as Anchors" \(PubMed\)](#) — The foundational first-mover experiments and the counter-anchoring remedy. Trust: high; lab-based.
- [Ames & Mason 2015, "Tandem Anchoring" \(Columbia PDF\)](#) — Source for the bolstering-range finding; 5 studies. The single most actionable counterintuitive result. Trust: high.
- [Mason, Lee, Wiley & Ames 2013, "Precise Offers Are Potent Anchors" \(Columbia PDF\)](#) — The pro-precision case. Trust: high, but read alongside the reversal below.
- [Loschelder et al. 2016, "The Too-Much-Precision Effect" \(PubMed\)](#) — The inverted-U reversal for experts; the crucial qualifier to "always be precise." Trust: high.
- [Hukkanen & Keloharju 2019, precision in real M&A \(SSRN\)](#) — Field data (\$4–5M round-number penalty). Trust: moderate-high; observational, so selection is a live confound.
- [Van Kleef, De Dreu & Manstead 2004, anger & happiness \(PDF\)](#) — The origin of "anger extracts concessions." Trust: high for the effect, but bounded by the meta below.
- [Sharma, Elfenbein, Sinha & Bottom 2020, emotion meta-analysis \(SSRN\)](#) — 64 studies; the definitive "anger is conditional" correction. Trust: high.
- [Galinsky, Maddux, Gilin & White 2008, perspective-taking vs. empathy \(PubMed\)](#) — Grounds the reframing of "tactical empathy." Trust: high.
- [Maddux, Mullen & Galinsky 2008, mimicry \(ScienceDirect\)](#) — The 67%-vs-12.5% agreement result; the closest lab cousin to Voss's mirroring — but a loose match. Trust: high for mimicry, low as Voss validation.
- [Mazei et al. 2015, gender meta-analysis \(APA PDF\)](#) — 123 effect sizes; the "small and moderated" correction to backlash. Trust: high.
- [Schaerer, Swaab & Galinsky 2015, "Anchors Weigh More Than Power" \(SAGE\)](#) — The weak-BATNA paradox. Trust: high; counterintuitive and well-designed.
- [Pew Research 2023, salary-negotiation outcomes](#) — The real distribution that debunks the "38% more" myth. Trust: high; survey data.
- [McKeown, critique of *Getting to Yes* \(SSRN PDF\)](#) — The scholarly case against win-win universalism. Trust: moderate; argument, not data.